

Agent based simulations and optimization with uniform grid applied to colloidal jamming transitions

Jimmy-Xuan Shen¹

Department of Physics, University of California, Santa Barbara, California 93106-9530, USA

I. INTRODUCTION

A. Agent based simulations

Agent based simulation (ABS) is a class of computational techniques used to simulate the behavior of autonomous individual agents with the intent to study the behavior of the collective. The technique is used to study a wide range of problems ranging from soft-matter physics and biology to social simulations and economics. A key feature of these simulations is that the interactions between the individual agents give rise to non-trivial behavior of the entire system. A common bottle-neck of ABS is iterating over all of the agents in a system and accessing the data stored for each agent to detect for interactions; for a system of N agents, there will be $\mathcal{O}(N^2)$ such interaction checks. Fortunately, in most cases it is only possible for a given agent to interact with a relatively small number of other agents at any given time. However, the iteration over all particles to check for interactions cannot be skipped using standard data storage techniques since we have no prior knowledge of the agent before accessing its information in the computer's memory. To avoid the unnecessary memory accesses, we will have to implement the memory storage of the agents in an organized way so that only a small set of possible interaction candidates will be accessed each time. In this report, a particular example of the ABS principle applied to colloidal jamming transitions will be discussed to provide a solid example. The example we are using is quite simple in terms of offering physical understanding of realistic colloids. However, the concepts can be used for any system where many individual elements interact over short distances.

B. Jamming transitions

The particular system we will be studying is a simple two-dimensional colloid consisting entirely of hard-disks. For simplicity the particles in the system will be treated as hard-disks with elastic collisions as the only interaction between them. We will also ignore all of the properties of the suspending liquid such as viscosity and temperature.

Jamming transitions^{1,2} are a physical process found in colloids and other complex materials like glass and foam, where the complex fluids become rigid with increasing density. While one might expect the transition from colloidal liquid to a jammed rigid state to occur at a well-defined density, recent studies have shown that this transition is not well defined but rather exists over a continuous range.³

For our simple system, we expect the particles to be able to move around freely, while the closely packed system will not be able to move at all. The system will evolve in discrete time dt , and the interaction checking problem in this case is the detection of collisions between all of the particles in the system. Speed-up of the simulation will come from knowing some information about the position of a particle before the stored data of that given particle is actually accessed.

The only physically relevant quantity is the filling fraction (ϕ) of total space occupied by the particles in the system — a unit-less parameter. Thus, we can use arbitrary units to describe the radial size (r) and velocity (v) of the particles in the system, with the requirement that $v \times dt$ is small enough such that possible collisions will not be skipped during any given time-step. For simplicity, the mass of all particles will be set to 1 in arbitrary units.

II. DESCRIPTION OF SIMULATION

A. Initialization

The particles are initialized in fixed positions in a bounded square box. They are also initialized with random velocities sampled from a uniform disk, with a maximum velocity of 1 (in reduced units). Because the number of particles in the system does not change during the course of a simulation, we can designate a continuous block of memory to store the physical properties of the particles (position and velocity) using a standard c++ array. The speed-up of our simulation is based on organizing the data of the particles into groups that are dependent on the position of the particles. The groups will be determined by the position of the centers of the particles on a uniform $k \times k$ grid, we will call each square created by the grid a “*neighborhood*”. Since the particles can move freely between the neighborhoods, we expect to constantly move elements between the groups. The most efficient data structure for frequent data relocations is a linked list,⁴ where each piece of data stored in the list form a “*node*” and all of the nodes are connected by memory references (or pointers) to other nodes. The pointers are variables in c++, like integers or floats, that take on the value of physical memory addresses in the computer. The addition/deletion of a node simply requires changing a few of the pointers in the list to accommodate for the new/removed node, Figure 1.

Because the N particles in the system will always use the same locations in memory to store their physical attributes, we only need to organize their memory references into linked lists based on which neighborhood the particle is in. Doing this allows us to utilize the compilers optimized array data access for tasks such as writing the particle data to file, while using the linked list data structures to optimize collision detection.

B. Update particle positions

To update the positions of the particles, we will move the particles one by one at each time-step and check for collisions with the other particles. To better simulate the geometry of each collision,




Figure 1. (a) Basic structure of a linked list. The arrows indicate pointer references to where the other nodes are stored in memory. (b) Reorganization of the pointers to insert an additional node to the list. (c) Reorganization of the pointers to delete a node to the list.

we will advance each particle over five sub time-steps, moving by $v \times \frac{dt}{5}$ a total of five times. Normally, we will have to check for collision of the advancing particle with all of the other $N - 1$ particles in the system, which will require $N - 1$ distance calculations. However, we can take advantage of our storage structure and only access a subset of the remaining particles in the system for possible collision candidates. Since the advancing particle cannot travel too far from the initial position in dt , we will restrict ourselves to a 3×3 block of neighborhoods (Figure 2) to look for colliding particles. Thus, instead of checking $N - 1$ particles for collision, we will be checking approximately $\frac{9}{k^2}(N - 1)$ particles, assuming a uniform distribution. If a collision has been found, meaning that advancing the particle by $v \times dt/5$ will result in overlap with another particle, the advancing particle is halted (skipping the remainder of the sub time-steps) and the velocities of both the advancing particle and the colliding particles are modified according to elastic collision of hard disks.

At the end of each time-step, we check to see if the particle has moved to a new neighborhood. If the particle is in a new neighborhood, we will




Figure 2. For a given particle (red), only collisions with the yellow particles will be considered possible in our simulation. The data of the green particles are not accessed during collision detection.

update the linked lists for the old and new neighborhoods accordingly. Note that doing this in a linked list only requires us to modify the pointers of at most two other elements in the list and will not disturb the rest of the data. Doing the same operation with an array will sometimes require the entire array to be rewritten.

C. Treatment of collision

To treat the elastic collisions, we first shift the velocities of the incident particles to the center of mass (CM) frame so that the incident particle will have equal but opposite momentums, and the same is true of the particles after collision, Figure 3 (a). In addition, since the kinetic energy must be conserved, the magnitudes of the initial and final velocities are all equal. The only remaining unknown is the relative angle θ between the incident and outgoing particles. To determine this angle, we must use the fact that the impulse of the collision is entirely along the impact vector \mathbf{b} —displacement vector between the two colliding particles. Let α be the angle between the incident particle velocity \mathbf{u}_1 and the impact vector \mathbf{b} . Since the resulting velocity \mathbf{v}_1 must have the same magnitude as \mathbf{u}_1 and the component perpendicular to the impact vector cannot change (assuming frictionless collisions), we must have that $\alpha + \theta/2 = \pi/2$, as shown in Figure 3 (b). Once we have θ , we can solve for \mathbf{v}_1 and \mathbf{v}_2 in the CM frame and translate them back to the rest




Figure 3. (a) Collision in the CM frame. The initial velocities \mathbf{u}_1 and \mathbf{u}_2 and the final velocities \mathbf{v}_1 and \mathbf{v}_2 all have the same magnitude. (b) \mathbf{u}_1 and \mathbf{v}_1 must form an isosceles triangle with angle θ between them, and α is the angle between \mathbf{u}_1 and \mathbf{b} .

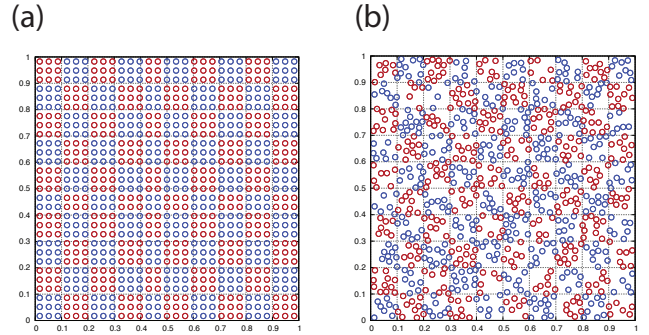


Figure 4. (a) Initial configuration of the particles. The uniform grid is shown and the particles in adjacent neighborhoods are colored differently. (b) The configuration of the particles after 100 time-steps.

frame.

The resulting snapshots of our simulations with $k = 10$ and $N = 841$ is shown in Figure 4.

III. DISCUSSION

A. Grid size dependence

To test the effectiveness of the uniform grid method, we can test the speed of the simulation using different grid sizes. For each grid size, we run a series of simulations with constant particle size and velocity distribution, but different number of particles. The time taken for each simulation run to complete 100 time-steps is shown in Figure 5. For each grid size, the systems with




Figure 5. Plot of time taken to complete 100 time-steps versus the number of particles in the simulated system, using 10×10 , 15×15 and 20×20 uniform grids. The two red circles have centers of the same y -value (Time), so while the system size more than doubled, there was little additional computation time.

more particles take more time to simulate. However, the relationship is not quite quadratic since particles in the dense system are not likely to complete all five sub time-steps before encountering another particle and halting. A key thing to note is that the simulation of 1000 particles using 10×10 grid takes the same amount of time as a 2000 particle simulation on a 20×20 grid. This is due to the fact that the total number of collision checks is $\mathcal{O}(9N^2/k^2)$, so scaling k and N together will not increase computational time. This is the primary benefit of the uniform grid method, as more particles can be included in the simulation with no additional computational cost (only need additional storage).

B. Where is the Jamming transition?

To quantify the mobility of the particles in our simulations, we store the initial positions of all particles and calculate the average distance of the particles to their initial positions. As we increase the density in our simulations, the particles are less likely to move freely without collision. And in a completely packed situation, the particles will not move at all. When we observe the average displacement of the particles in the system as a function of time (Figure 6), we see that the simu-

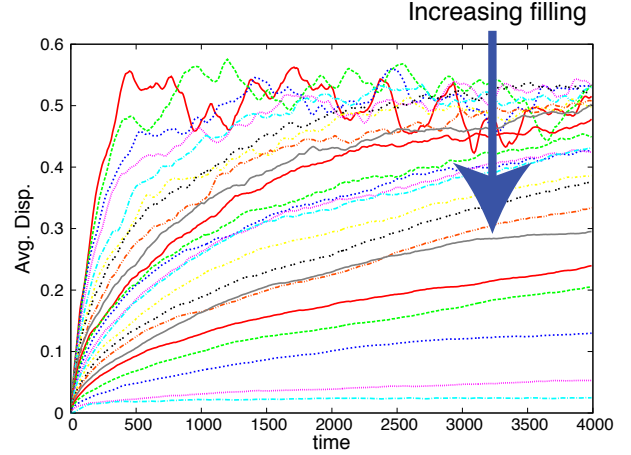


Figure 6. Time-series of the average displacement (as a fraction of the simulation system size) of particles in the system. All average displacement curves decrease as filling fraction increases.

lation exhibits some degree of mobility at almost all particle densities. This agrees with previous studies³ which found that the Jamming transition occurs over a continuous range and we cannot define an exact transition point. Although they modified the initial conditions by starting at different volumes and compressing the systems to a jamming state. The fact that our simulation shows some degree of sustained mobility for almost all particle densities is due to us treating all collisions as purely elastic. Our result cannot be directly compared with previously studies, which used a potential for inter-particle interactions and did not assume purely hard-sphere collisions. In Figure 7 we plot the average displacement after 100,000 time-steps against the filling fractions of the systems, for both square and triangular lattice of initial particle positions. The triangular lattices exhibit lower mobility at high filling fractions because each particle in the triangular initial arrangement has *six* nearest neighbors (compared to *four* for the square lattice). This means the triangular lattice will have less gaps between the nearest neighbors for a given particle to move at each time step. At lower filling fractions both arrangements gives average displacement of 0.5, as the final positions are roughly randomly distributed.

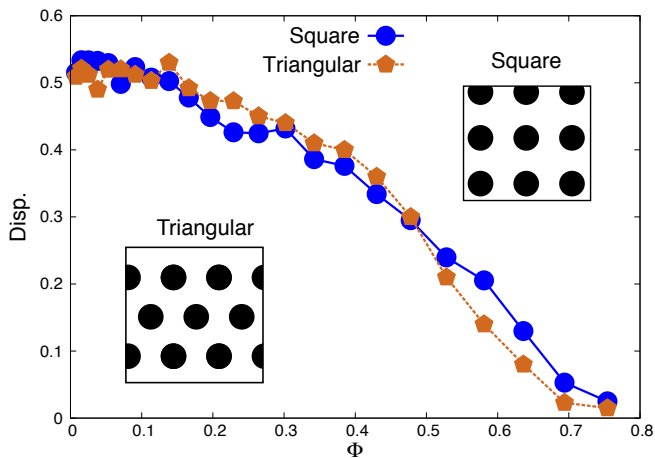


Figure 7. Plot of the average displacement of particles in the system against the filling fraction for both square and triangular initial particle arrangements. Insets show illustrations of the initial configurations.

C. Additional applications

Since this study is focused on the implementation of the uniform grid method, it is worthwhile to discuss other potential application of the method. One important aspect of linked list that we have not fully utilized is the fact that an addition or deletion of nodes in the list does not require any kind of array resizing (which require many memory operations). This is important in simulating systems where the total number of agents is not conserved; for example, in reaction of molecules adsorbing on the surface of a catalyst. For these types of systems, we will not be able to store all of the agents in a fixed array as we have done for the jamming transition. We will need to store the full data of each agent in a linked list instead of storing just the pointer to its memory

address. This allows for not only frequent relocation of data between the linked lists, but also for permanent addition and deletion of data of the whole system, which frees up memory so that new particles can be created later (avoiding memory leaking in long simulations).

IV. CONCLUSION

Using the simple example of hard sphere collisions, we have demonstrated the power and effectiveness of the uniform grid method, which is able to scale up the size of the simulation without additional CPU time. Using our simulations we found that the jamming transition is difficult to identify due to the fact that jamming depends on the initial configuration of the system. This agrees in spirit with previous theoretical studies. A more practical definition of where the jamming transition occurs could be where the average displacement of the particles is below 10% the size of the system after a sufficiently long time.

REFERENCES

- ¹C. O'Hern, L. Silbert, A. Liu, and S. Nagel, Phys. Rev. E **68**, 011306 (2003).
- ²G. Biroli, Nat Phys **3**, 222 (2007).
- ³M. Ozawa, T. Kuroiwa, A. Ikeda, and K. Miyazaki, Phys. Rev. Lett. **109**, 205701 (2012).
- ⁴J. L. Antonakos and K. C. Mansfield, *Practical Data Structures Using C/C++*, 1st ed. (Prentice Hall, Upper Saddle River, N.J, 1998).